DIGITAL METHODS FOR ANALYSING TEXTS
//

02_Analysing text

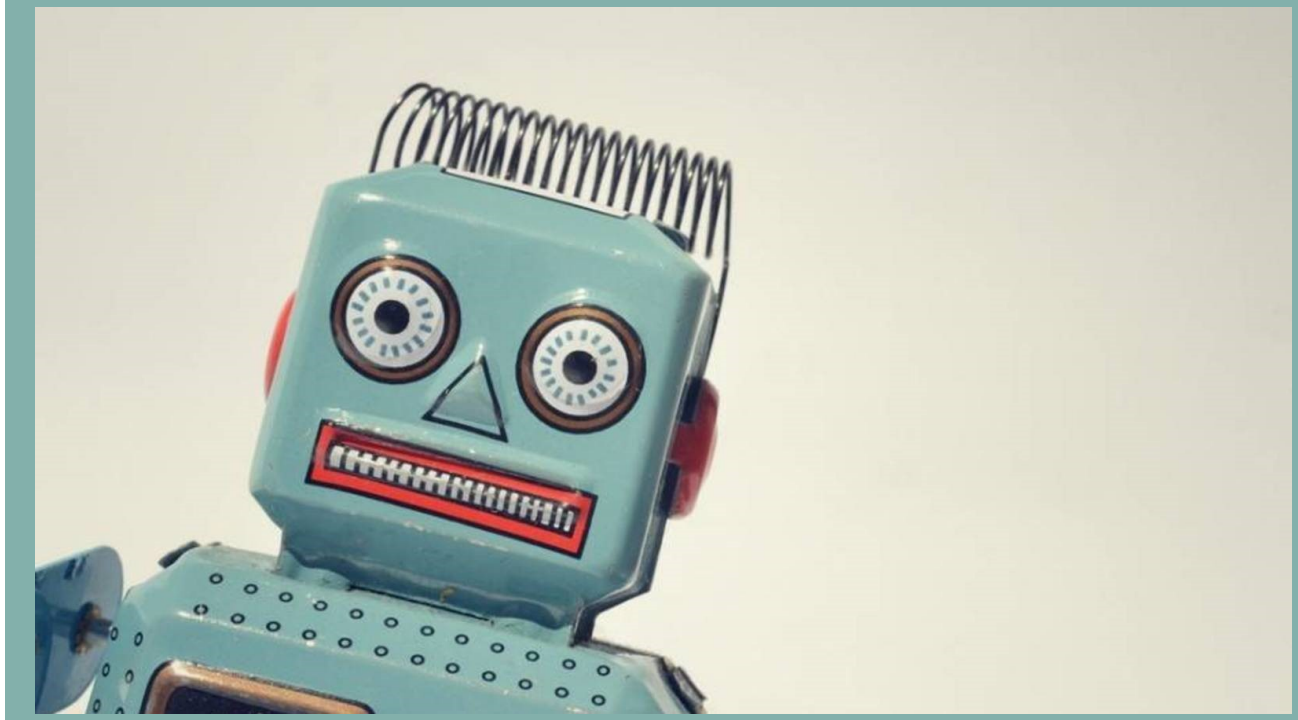

Ana Valdivia
Research Associate
King's College London

BEFORE WE START...//

What have we learned?

- Definition of AI and NLP.
- History of NLP.
- Language as data.

- Import unannotated and annotated corpus in Jupyter Notebook (Python).

BEFORE WE START…//

Difference between words and tokens

- If a text is 1,000 **words** long, it is said to have 1,000 **tokens**.

- But a lot of these words will be repeated, and there may be only say **400 different words** (**types**) in the texts.

ROAD MAP//

1. Corpus Preprocessing

2. Representing text

- 2.1. Enter the matrix
- 2.2. Discrete Representations
- 2.3. Distributed Representations
- 5.4. Discrete versus Continuous

CORPUS PREPROCESSING //

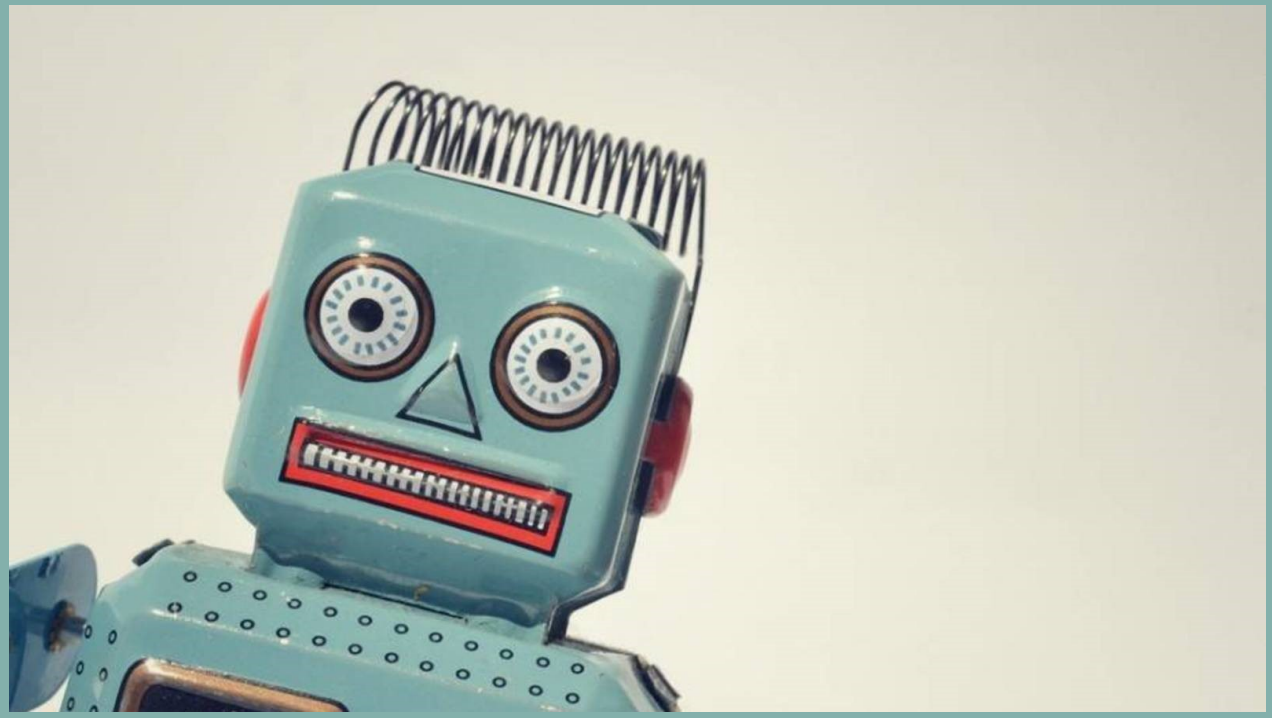

CORPUS PREPROCESSING//

How would you manually represent a corpus?

Click here

CORPUS PREPROCESSING//

Think like a person would do it,
and translate it into code.

CORPUS PREPROCESSING//

Office documents

??????

Features:

Term-matrix documents

Embeddings

Docs	amp	brexit	euref	leav	remain	strongerin	vote	voteleav
738102860454498304	2	1	1	0	0	0	1	1
739933062281187329	0	0	1	2	2	0	1	0
745289444006170624	0	0	0	1	1	0	4	0
745501761289355264	0	0	0	0	7	0	0	0
745621915516149760	0	1	1	1	1	0	2	0
745649059231215616	1	0	0	1	1	1	2	0
745875415839965184	2	0	1	0	1	0	2	0
745922585494429697	1	0	1	0	1	1	2	0
745973624142725120	2	0	0	1	1	1	1	0
746108821479821312	0	0	1	0	4	0	1	0

CORPUS PREPROCESSING//

Document decomposition

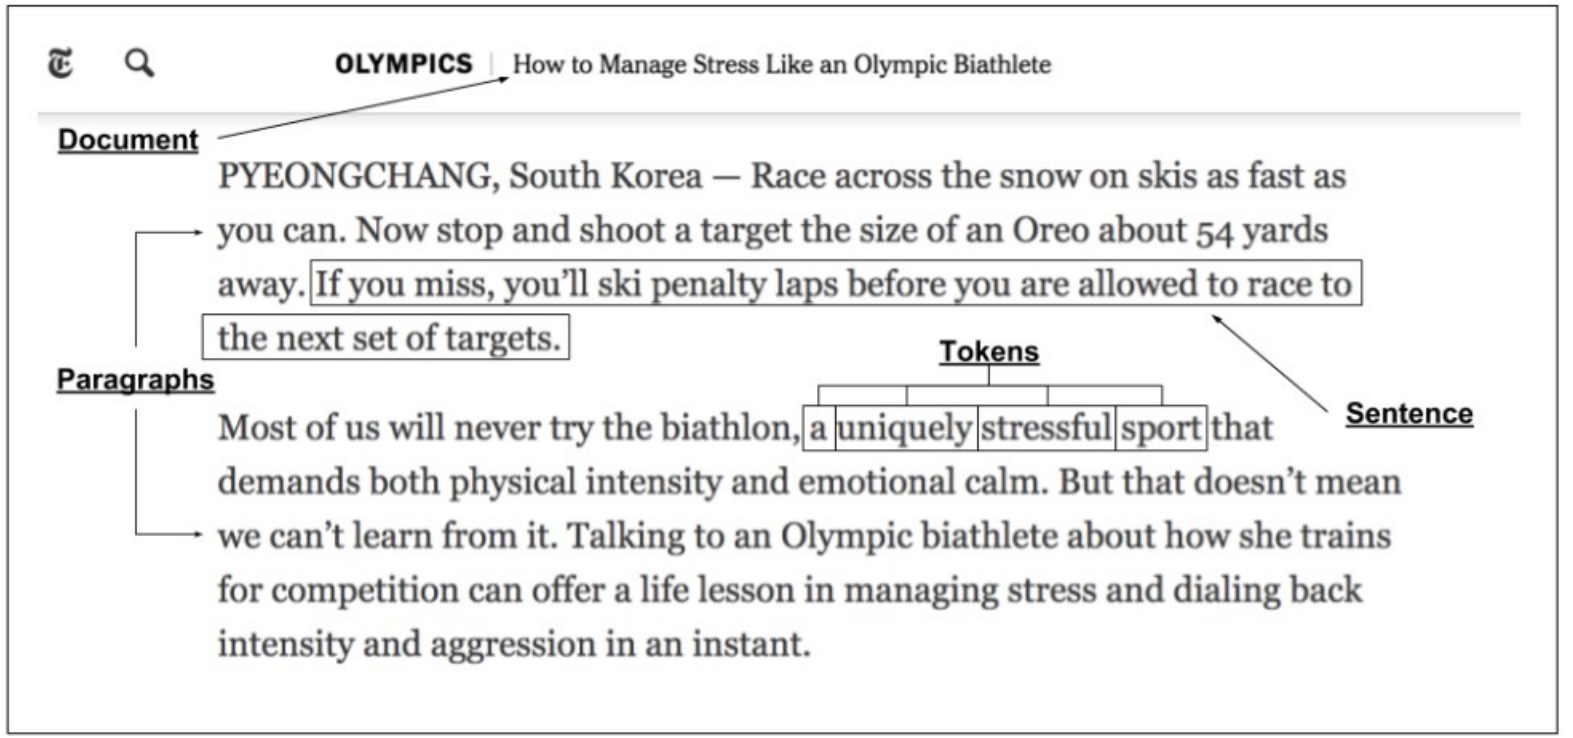

Figure 3-2. Document decomposition illustrating the distribution of meaning across paragraphs, sentences, and individual tokens

Features

Discrete features

Represents a feature with a specific meaning.

Term-document matrix.

Continuous features

Do not *mean* anything anymore.

They can not be interpreted.

Word embeddings.

Unicode

Our programs will often need to deal with different languages, and different character sets.

- If you live in the English-speaking world you probably use ASCIIT.
- If you live in Europe you might use one of the extended Latin character sets, containing such characters as:
  - “ø” for Danish and Norwegian,
  - “ő” for Hungarian,
  - “ñ” for Spanish and Breton, and
  - “ň” for Czech and Slovak.

CORPUS PREPROCESSING//

Unicode

Unicode supports over a million characters. Each character is assigned a number, called a code point.

In Python, code points are written in the form \uXXXX, where XXXX is the number in four-digit hexadecimal form.

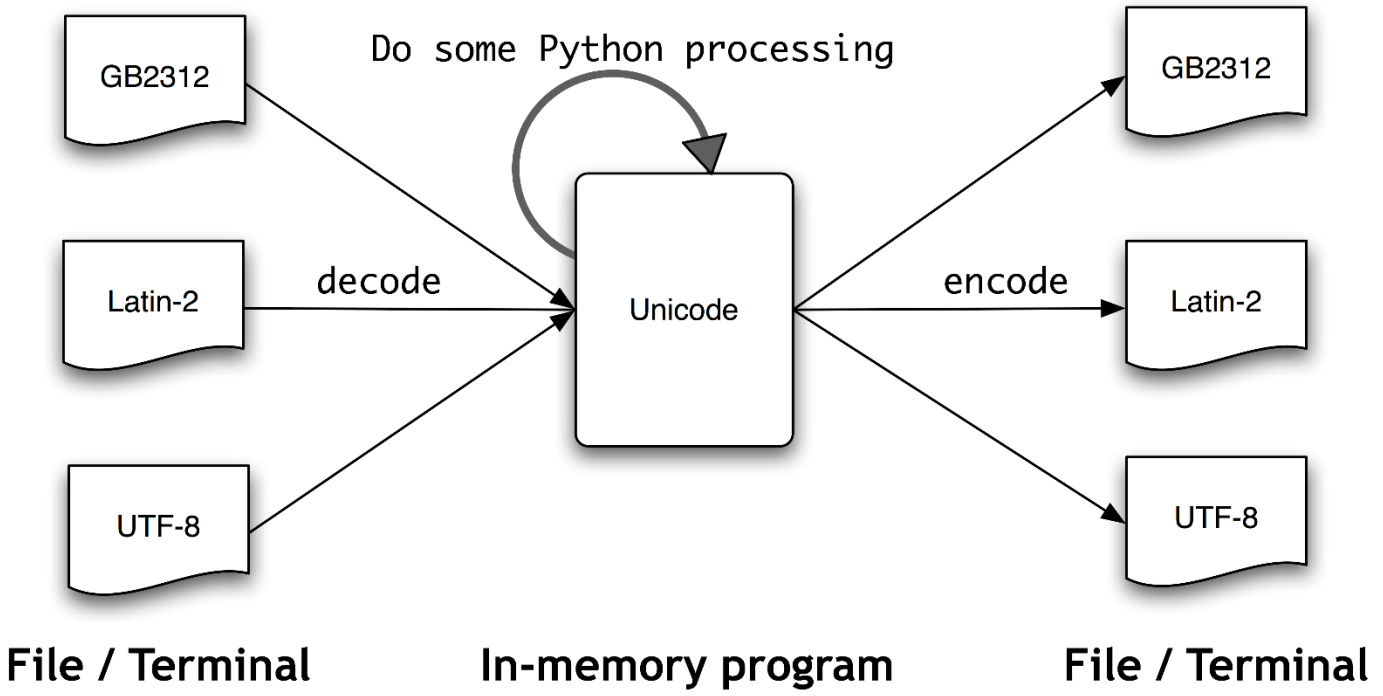

Regular Expressions (regexpr)

Many linguistic processing tasks involve pattern matching. For example, we can find words ending with `ed` using `endswith('ed')`.

Regular expressions give us a more powerful and flexible method for describing the character patterns we are interested in.

Table 3-3. Basic regular expression metacharacters, including wildcards, ranges, and closures

Operator	Behavior
`.`	Wildcard, matches any character
`^abc`	Matches some pattern *abc* at the start of a string
`abc$`	Matches some pattern *abc* at the end of a string
`[abc]`	Matches one of a set of characters
`[A-Z0-9]`	Matches one of a range of characters
`ed|ing|s`	Matches one of the specified strings (disjunction)
`*`	Zero or more of previous item, e.g., `a*`, `[a-z]*` (also known as *Kleene Closure*)
`+`	One or more of previous item, e.g., `a+`, `[a-z]+`
`?`	Zero or one of the previous item (i.e., optional), e.g., `a?`, `[a-z]?`
`{n}`	Exactly *n* repeats where *n* is a non-negative integer
`{n,}`	At least *n* repeats
`{,n}`	No more than *n* repeats
`{m,n}`	At least *m* and no more than *n* repeats
`a(b|c)+`	Parentheses that indicate the scope of the operators

REPRESENTING TEXT
//

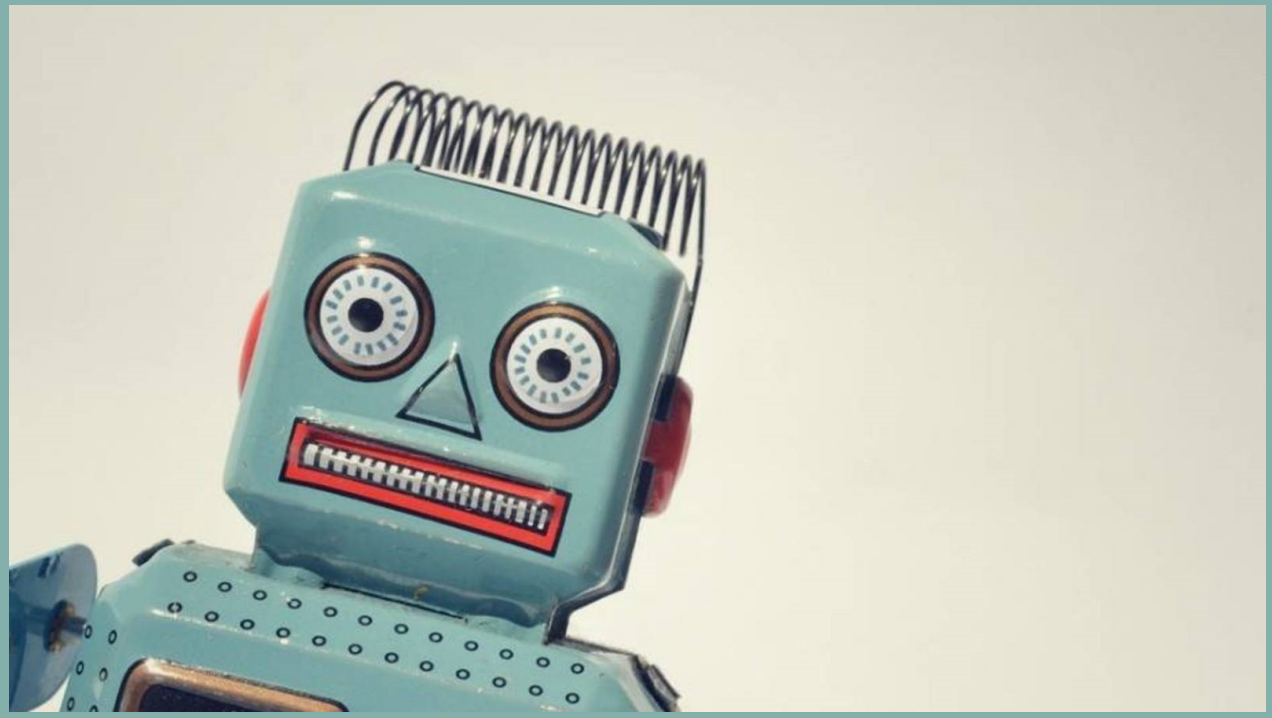

REPRESENTING TEXT//

Discrete features

N-gran features

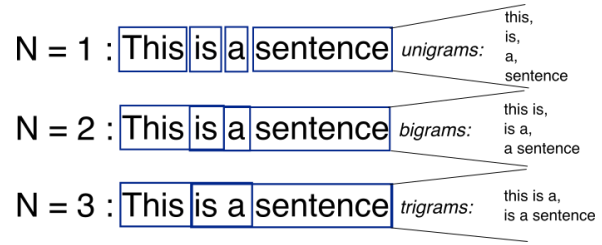

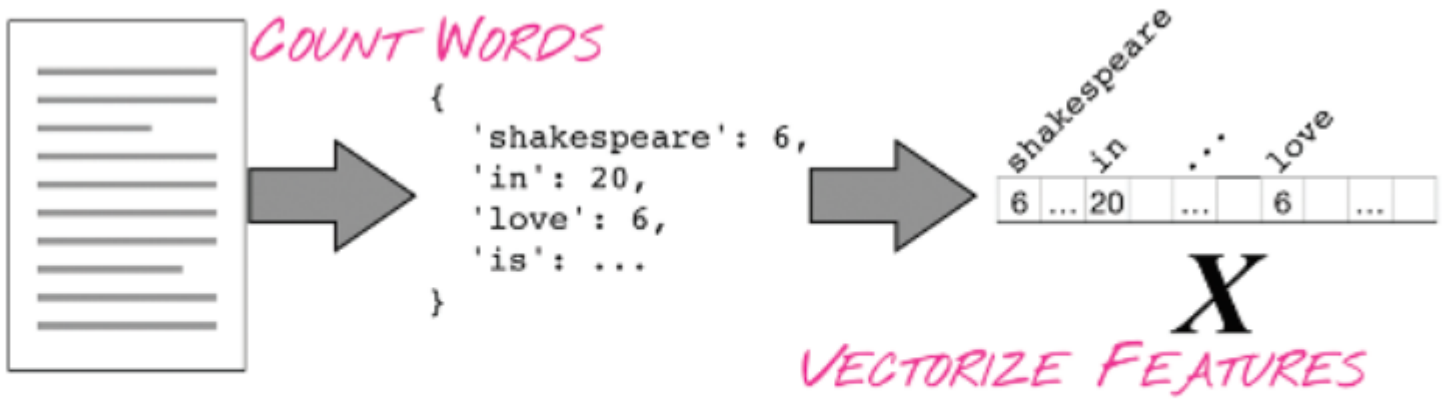

Figure 4 Schematic of a bag-of-words representation.

REPRESENTING TEXT//

The Term-Document matrix

A **document-term matrix** is a mathematical matrix that describes the frequency of terms that occur in a collection of documents. In a document-term matrix, rows correspond to documents in the collection and columns correspond to terms.

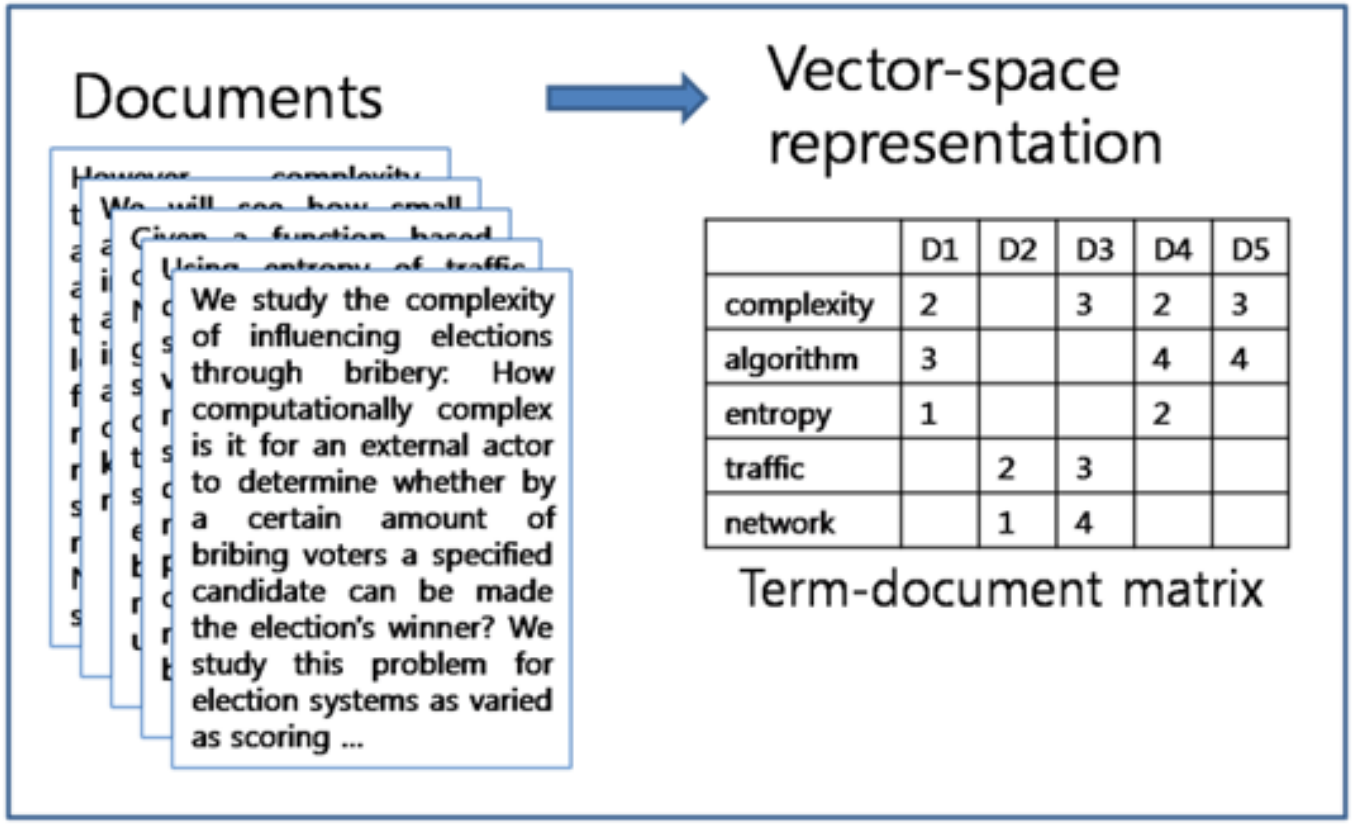

	D1	D2	D3	D4	D5
complexity	2		3	2	3
algorithm	3			4	4
entropy	1			2	
traffic		2	3		
network		1	4		

Term-document matrix

REPRESENTING TEXT//

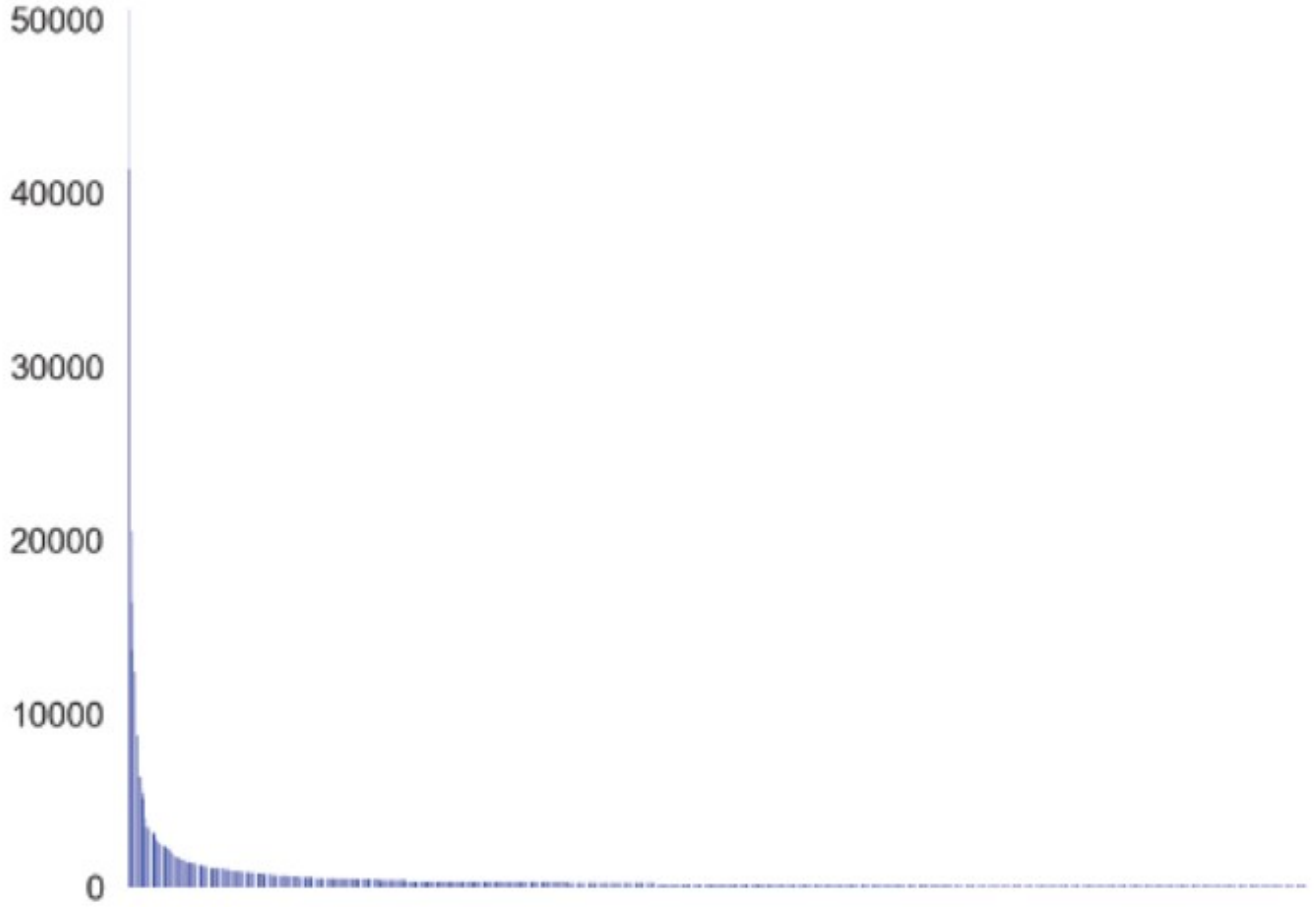

Figure 3 Frequency distribution of the top 1,000 words in a random sample of tweets, following Zipf's law.

REPRESENTING TEXT//

TF-IDF Counts

The **term frequency–inverse document frequency (tf-idf)**, is a numerical statistic that is intended to reflect how important a word is to a document in a collection or corpus.

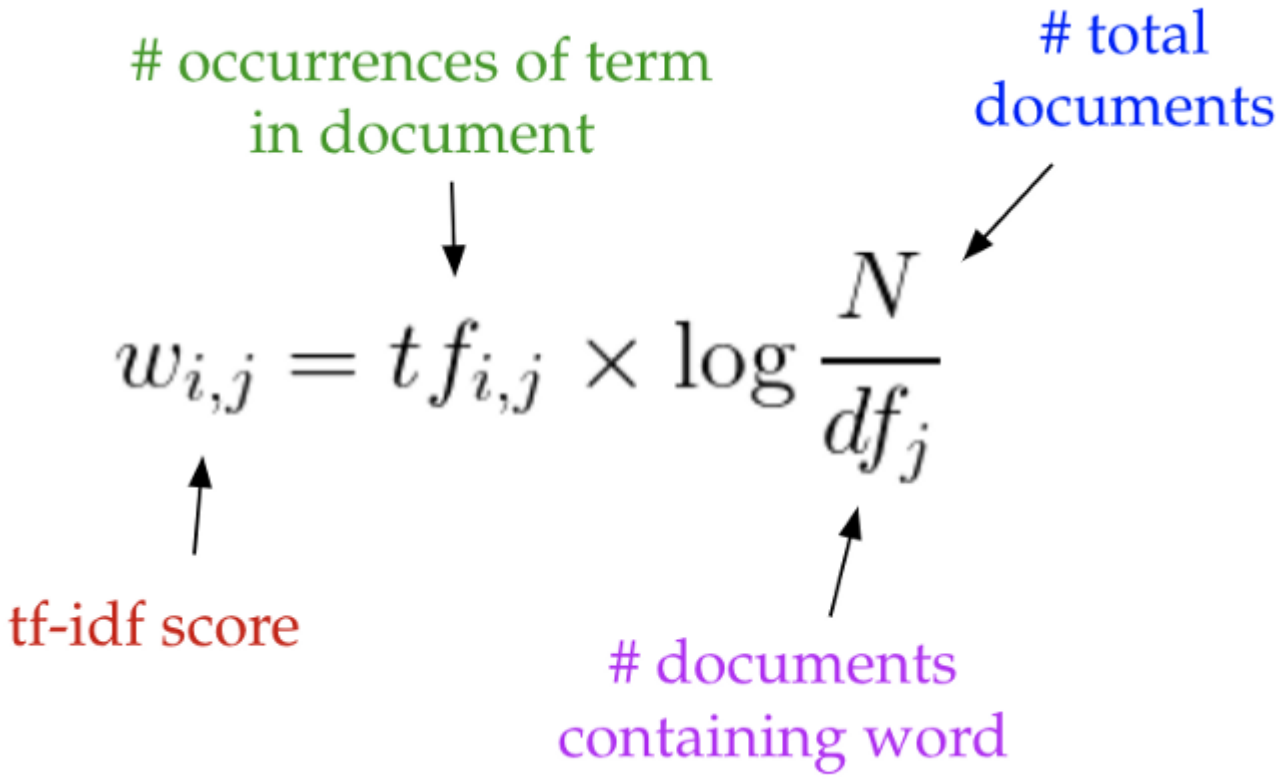

LET'S CODE!

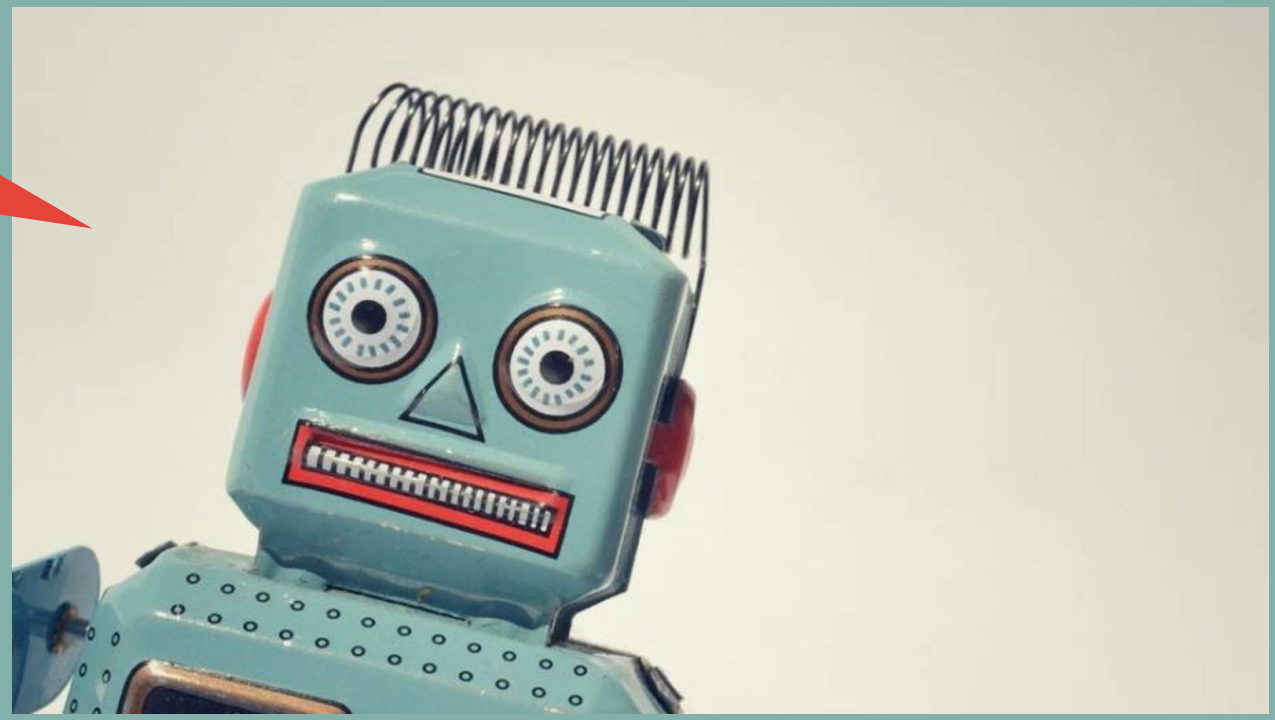

WE'LL BE BACK IN 15 MIN...